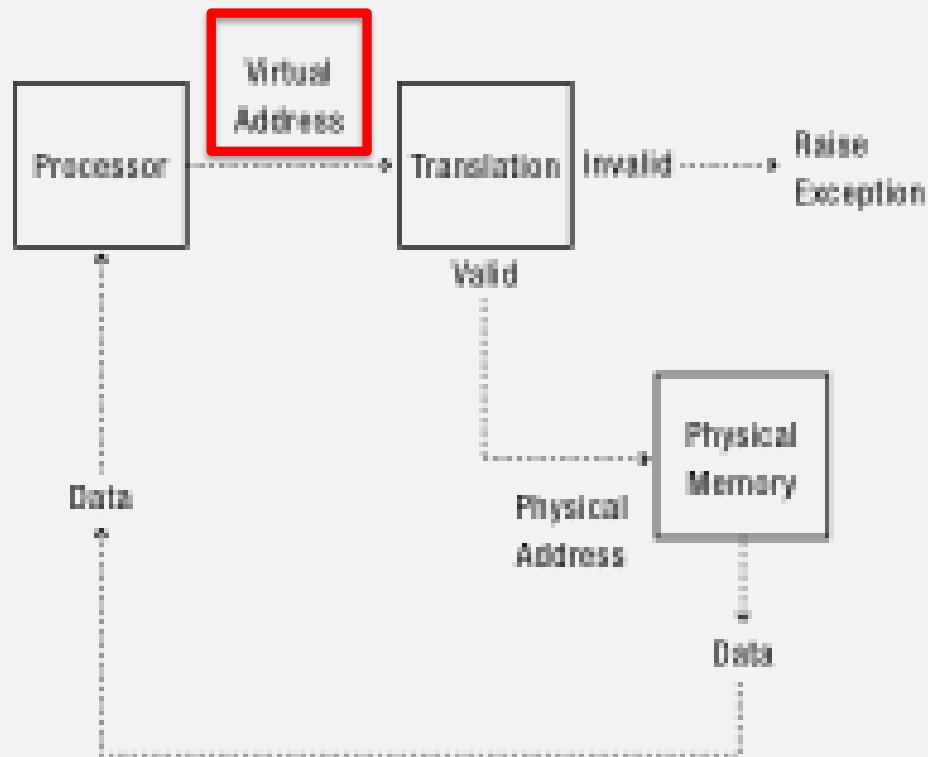


Address Translation

Main Points

- Address Translation Concept
 - How do we convert a virtual address to a physical address?
- Flexible Address Translation
 - Base and bound
 - Segmentation
 - Paging
 - Multilevel translation
- Efficient Address Translation
 - Translation Lookaside Buffers
 - Virtually and physically addressed caches

Address Translation Concept



Address Translation Goals


- Memory protection
- Memory sharing
 - Shared libraries, interprocess communication
- Sparse addresses
 - Multiple regions of dynamic allocation (heaps/stacks)
- Efficiency
 - Memory placement
 - Runtime lookup
 - Compact translation tables

Bonus Feature

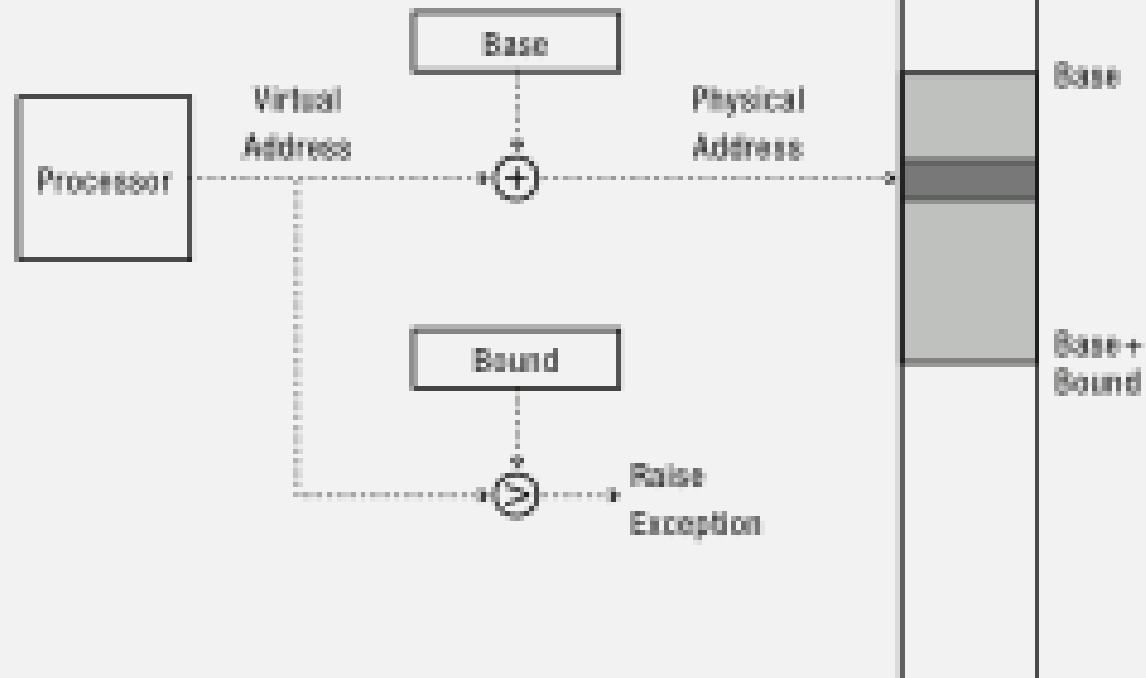
- What can you do if you can (selectively) gain control whenever a program reads or writes a particular virtual memory location?
- Examples:
 - Copy on write
 - Fill on demand
 - Demand paging
 - Memory mapped files
 - ...

Virtually Addressed Base and Bounds

Processor's View



Implementation



Activity #1

- Drawback of Base and Bounds

Virtually Addressed Base and Bounds

• **Virtual Addressing** is a technique for mapping virtual addresses to physical addresses. It is used in operating systems to manage memory and protect processes from each other. Virtual addressing allows for the use of a virtual address space that is larger than the physical address space. The virtual address space is divided into virtual pages, which are mapped to physical pages. The mapping is done by the operating system using a virtual-to-physical address translation table. This table is stored in memory and is updated as the process runs. Virtual addressing is a key component of memory management in operating systems.

• **Base and Bounds** are two techniques for implementing virtual addressing. The base technique uses a base register to store the starting address of the virtual address space. The bounds technique uses a bounds register to store the ending address of the virtual address space. Both techniques use a virtual-to-physical address translation table to map virtual addresses to physical addresses. The base technique is simpler to implement but is less flexible. The bounds technique is more flexible but is more complex to implement. Both techniques are used in operating systems to manage memory and protect processes from each other.

• **Virtual Addressing** is a technique for mapping virtual addresses to physical addresses. It is used in operating systems to manage memory and protect processes from each other. Virtual addressing allows for the use of a virtual address space that is larger than the physical address space. The virtual address space is divided into virtual pages, which are mapped to physical pages. The mapping is done by the operating system using a virtual-to-physical address translation table. This table is stored in memory and is updated as the process runs. Virtual addressing is a key component of memory management in operating systems.


• **Base and Bounds** are two techniques for implementing virtual addressing. The base technique uses a base register to store the starting address of the virtual address space. The bounds technique uses a bounds register to store the ending address of the virtual address space. Both techniques use a virtual-to-physical address translation table to map virtual addresses to physical addresses. The base technique is simpler to implement but is less flexible. The bounds technique is more flexible but is more complex to implement. Both techniques are used in operating systems to manage memory and protect processes from each other.

Segmentation

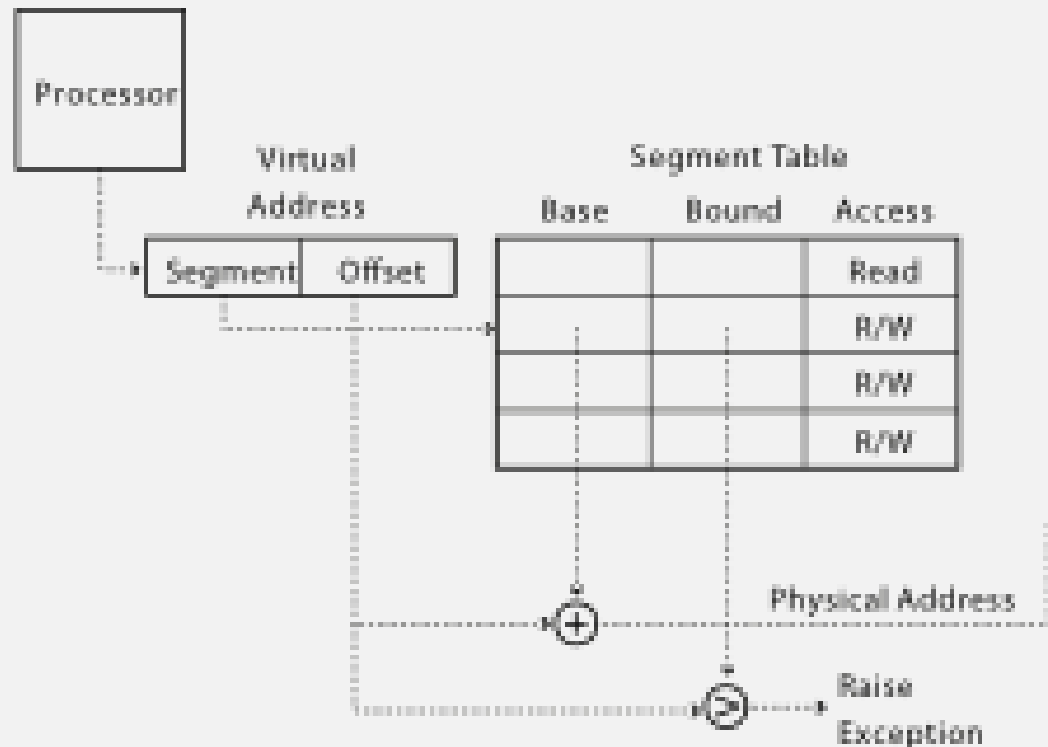
- Segment is a contiguous region of *virtual* memory
- Each process has a segment table (in hardware)
 - Entry in table = segment
- Segment can be located anywhere in physical memory
 - Each segment has: start, length, access permission
- Processes can share segments
 - Same start, length, same/different access permissions

Segmentation


Processor's View



Implementation



Physical Memory




	Segment	start	length	
2 bit segment #	code	0x4000	0x700	
12 bit offset	data	0	0x500	
	heap	–	–	
Virtual Memory	stack	0x2000	0x1000	Physical Memory


main: 240	store #1108, r2	x: 108	a b c \0
244	store pc+8, r31	...	
248	jump 360	main: 4240	store #1108, r2
24c		4244	store pc+8, r31
...		4248	jump 360
strlen: 360	loadbyte (r2), r3	424c	
...
420	jump (r31)	strlen: 4360	loadbyte (r2), r3
...		...	
x: 1108	a b c \0	4420	jump (r31)
...		...	

Processor's View

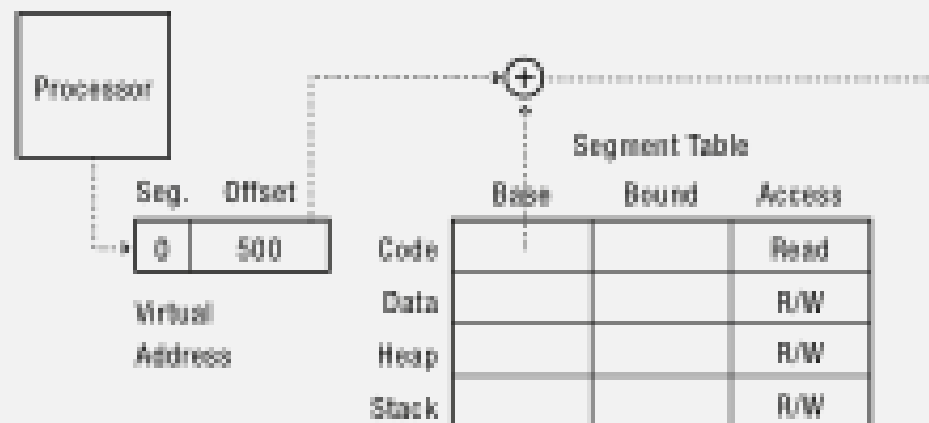

Process 1's View



Process 2's View



Implementation



Physical Memory



Activity #2


- Drawback of Segmentation

Segmentation

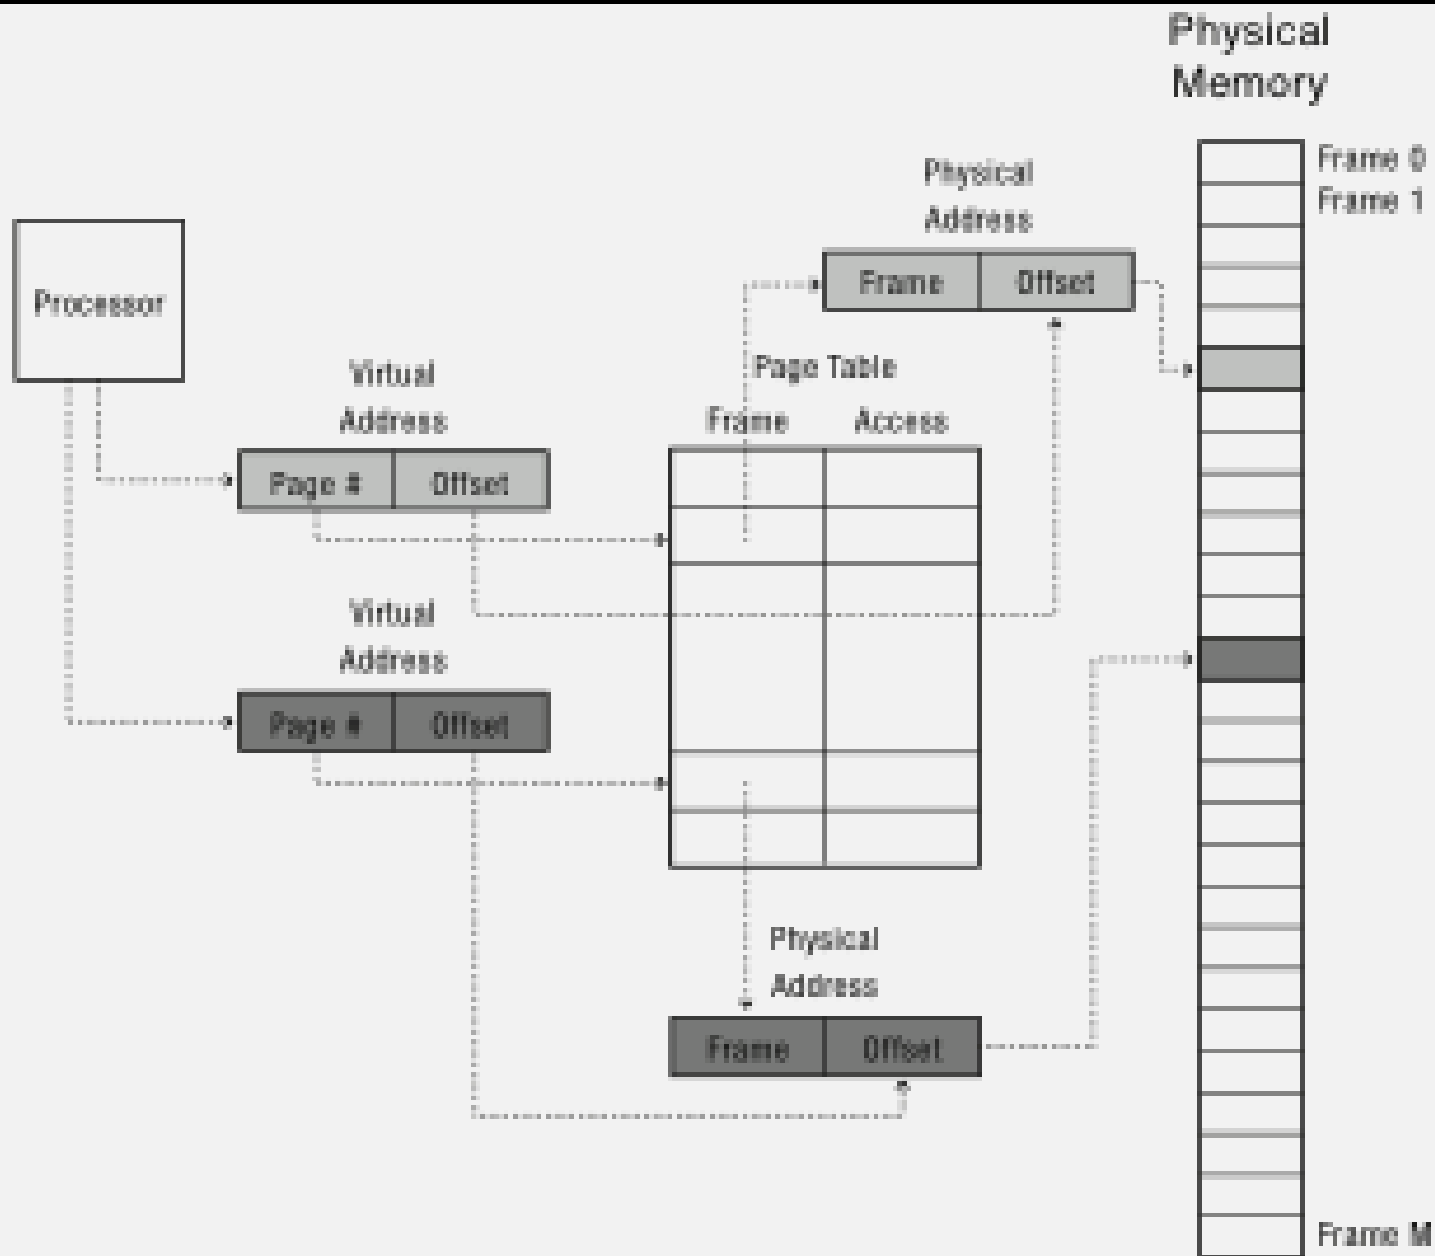
Paged Translation

- Manage memory in fixed size units, or pages
- Finding a free page is easy
 - Bitmap allocation: 00111111000000001100
 - Each bit represents one physical page frame
- Each process has its own page table
 - Stored in physical memory
 - Hardware registers
 - pointer to page table start
 - page table length

Paged Translation (Abstract)



Paged Translation (Implementation)



Process View

A
B
C
D
E
F
G
H
I
J
K
L

Physical Memory

I
J
K
L
E
F
G
H
A
B
C
D

Page Table

4
3
1

Activity #3

- Drawback of paging

Sparse Address Spaces

- Might want many separate dynamic segments
 - Per-processor heaps
 - Per-thread stacks
 - Memory-mapped files
 - Dynamically linked libraries
- What if virtual address space is large?
 - 32-bits, 4KB pages => 500K page table entries
 - 64-bits => 4 quadrillion page table entries

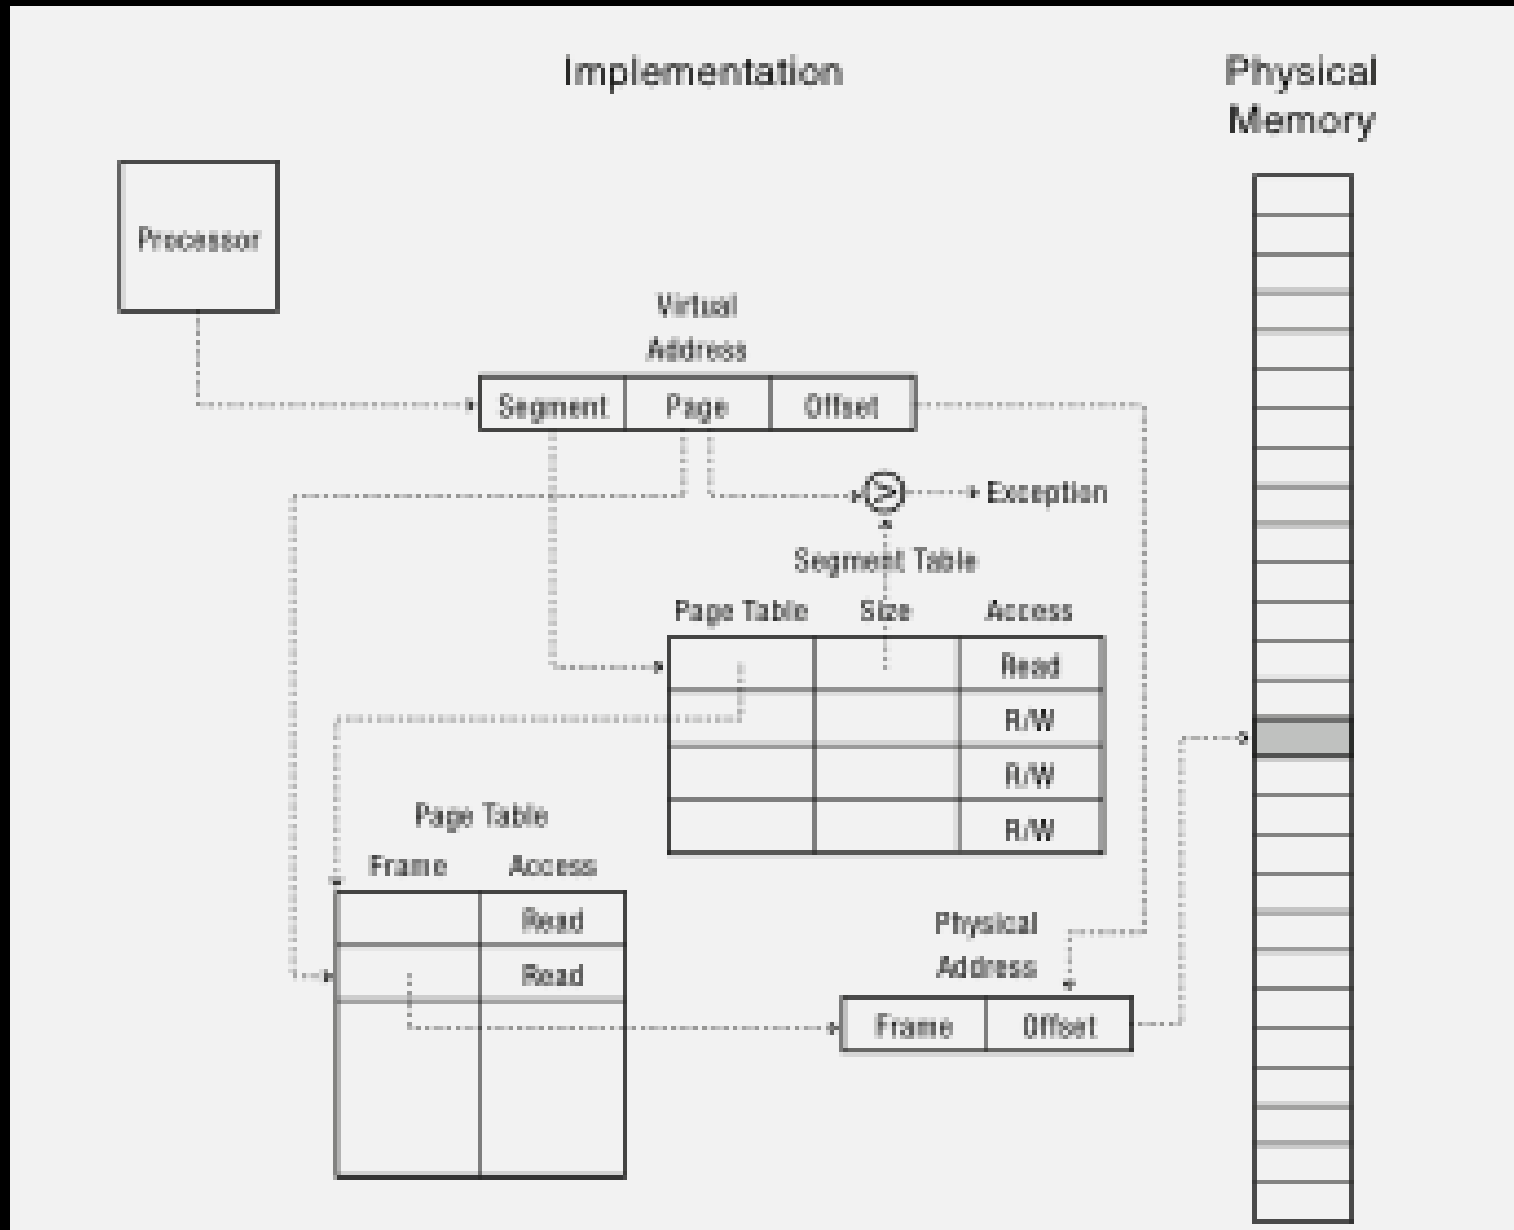
Multi-level Translation

- Tree of translation tables
 - Paged segmentation
 - Multi-level page tables
 - Multi-level paged segmentation
- Fixed-size page as lowest level unit of allocation
 - Efficient memory allocation (compared to segments)
 - Efficient for sparse addresses (compared to paging)
 - Efficient disk transfers (fixed size units)
 - Easier to build translation lookaside buffers
 - Efficient reverse lookup (from physical -> virtual)
 - Variable granularity for protection/sharing

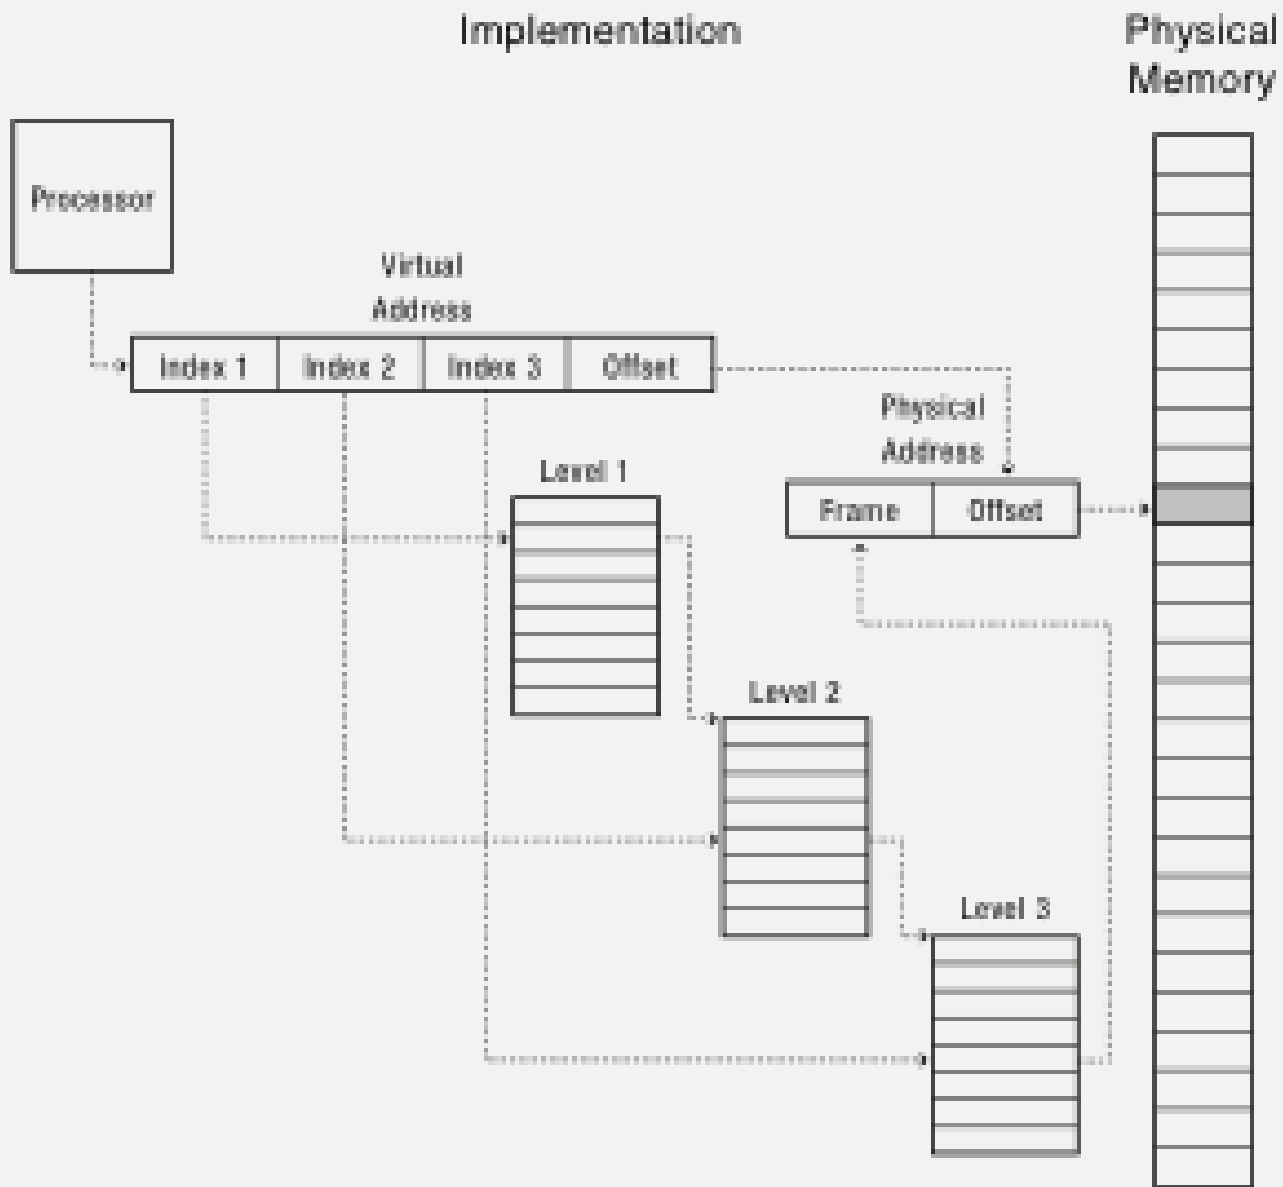
Paged Segmentation

- Process memory is segmented
- Segment table entry:
 - Pointer to page table
 - Page table length (# of pages in segment)
 - Access permissions
- Page table entry:
 - Page frame
 - Access permissions
- Share/protection at either page or segment-level

Paged Segmentation (Implementation)



Multilevel Paging



Activity #4

- Drawback of Multilevel translation

Multilevel Translation

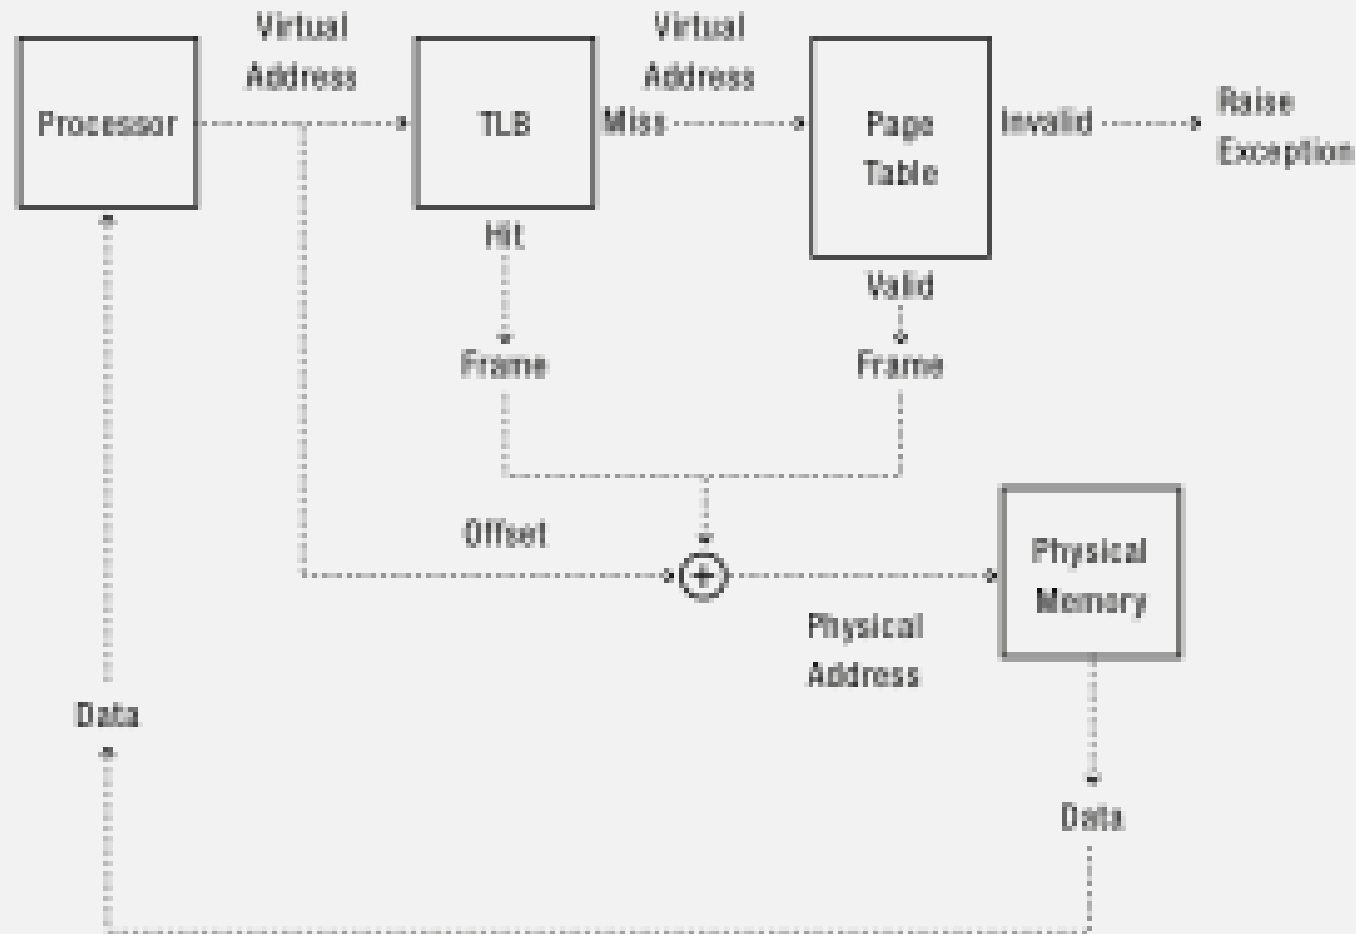
Activity #5

- Accelerate multilevel translation
 - Possible?
 - How?

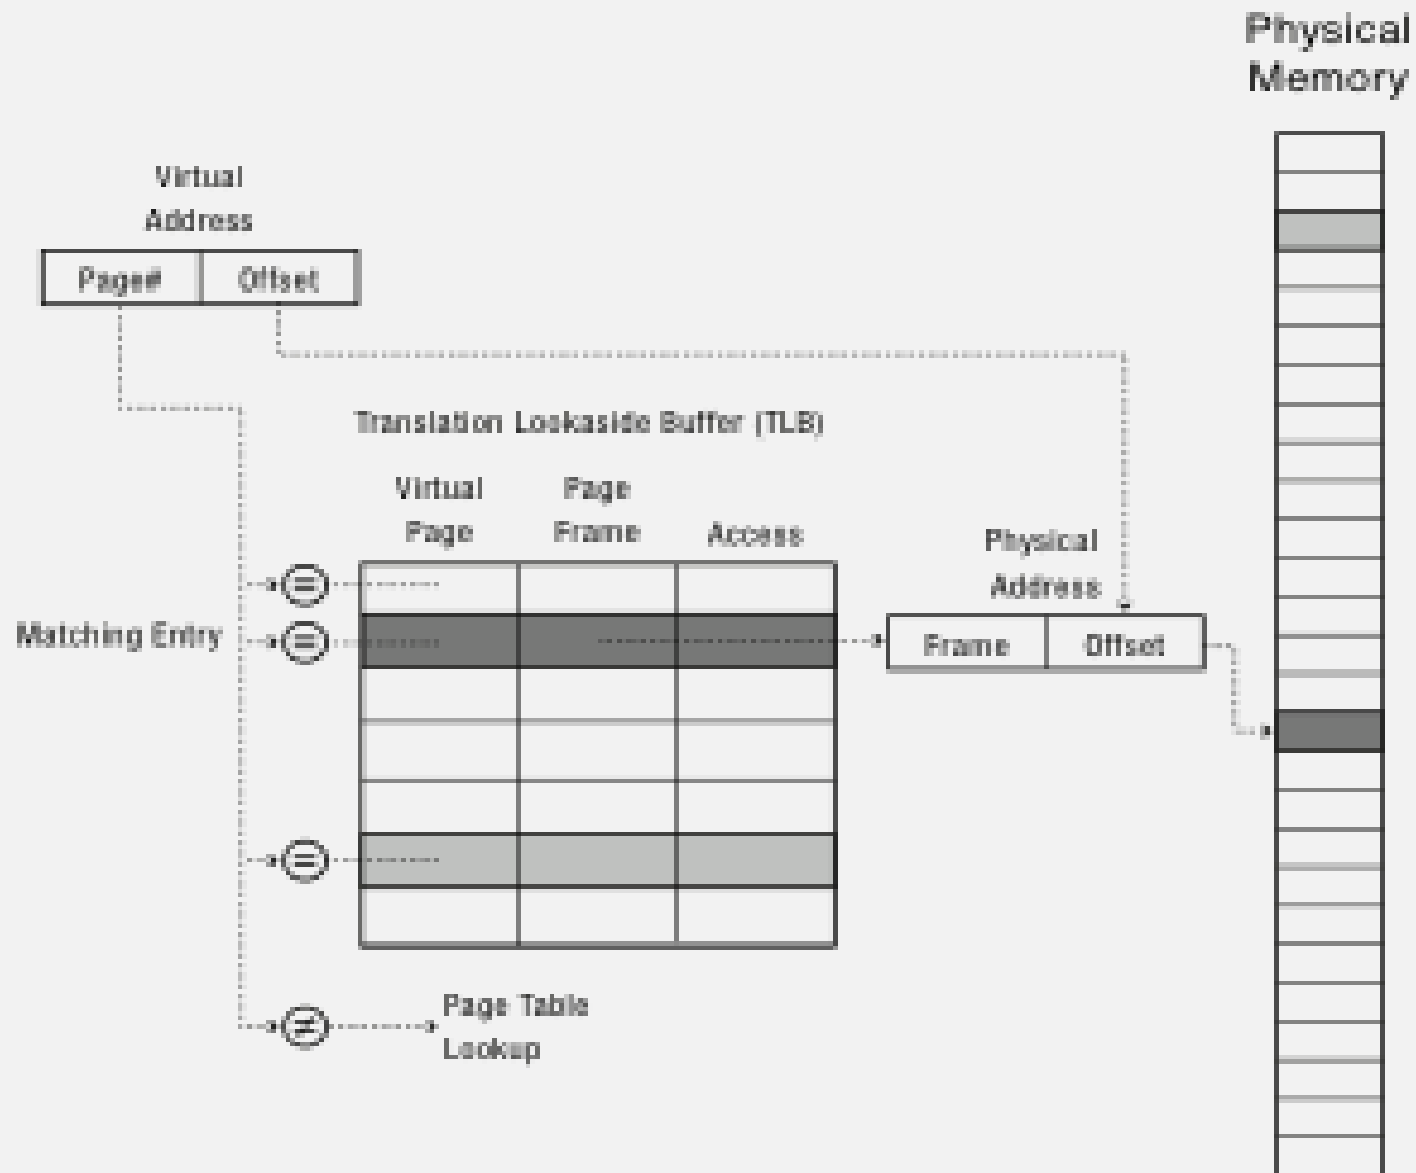
Efficient Address Translation

- Translation lookaside buffer (TLB)
 - Cache of recent virtual page \rightarrow physical page translations
 - If cache hit, use translation
 - If cache miss, walk multi-level page table
- Cost of translation =
Cost of TLB lookup +
 $\text{Prob}(\text{TLB miss}) * \text{cost of page table lookup}$

TLB and Page Table Translation



TLB Lookup



Address Translation Uses

- Process isolation
 - Keep a process from touching anyone else's memory, or the kernel's
- Efficient interprocess communication
 - Shared regions of memory between processes
- Shared code segments
 - E.g., common libraries used by many different programs
- Program initialization
 - Start running a program before it is entirely in memory
- Dynamic memory allocation
 - Allocate and initialize stack/heap pages on demand

Address Translation (more)

- Cache management
 - Page coloring
- Program debugging
 - Data breakpoints when address is accessed
- Zero-copy I/O
 - Directly from I/O device into/out of user memory
- Memory mapped files
 - Access file data using load/store instructions
- Demand-paged virtual memory
 - Illusion of near-infinite memory, backed by disk or memory on other machines